

Solutions to Problem Set 7

Data Compression With Deep Probabilistic Models

Prof. Robert Bamler, University of Tuebingen

Course material available at <https://robamler.github.io/teaching/compress21/>

Problem 7.1: Streaming ANS I — Implementation

The accompanying jupyter notebook contains our naive implementation of the Asymmetric Numeral Systems (ANS) algorithm. As we discussed at the end of the lecture, this implementation is not yet practically useful because the runtime cost for encoding k symbols scales quadratically in k . In this problem, you will implement the last missing piece of the ANS algorithm, which will reduce the runtime cost to just $O(k)$.

Problem. The reason why the current implementation suffers from $O(k^2)$ runtime cost is that the cost of each one of the k encoding or decoding operations is proportional to the amount of data that has already been compressed. Consider, e.g., the following line in the `decode` method:

```
self.compressed = self.compressed * scaled_probabilities[symbol] + z
```

Here, `self.compressed` is an integer whose binary expansion is the compressed representation of the symbols you’ve encoded so far. Thus, if you’ve already compressed, say, one megabyte worth of information content then `self.compressed` will be an eight-million-bit long integer. Multiplying it with `scaled_probabilities[symbol]` as in the above statement will, in general, change *all* of these eight million bits, which is expensive.

Fractional Bit Shifts. Figure 1 (i) provides a suggestion for a more intuitive interpretation of the encoding and decoding operations. Picture some buffer that stores the value of `self.compressed`. The allocated memory for this buffer will typically be a few bits larger than precisely necessary, and so the binary expansion of `self.compressed` will be padded with some leading zero bits (blue). The *valid bits* (i.e., everything except leading zeros) fit into a region of size $\log_2(\text{self.compressed})$ (up to rounding effects) that is *right-aligned* within the allocated memory (red box in Figure 1 (i)).

Let’s now look at the above line of code, which multiplies `self.compressed` with `scaled_probabilities[symbol]`. We will denote the latter as $m(x_i)$ from here on in parity with the lecture notes. If $m(x_i)$ is an integer power of two, i.e., if $m(x_i) = 2^r$ for some $r \in \mathbb{N}$ then the multiplication with $m(x_i)$ results in a simple left shift by $\log_2 m(x_i) = r$ bits (illustrated for $r = 3$ in Figure 1 (i)). Alternatively, we could simply append r zero bits at the end of the allocated memory, which would amount to only a constant run-time cost.¹

¹Technically, the *amortized* cost is constant, provided that we use a dynamic array (aka a “vector”).




Figure 1: (i) Multiplying an integer by a power of two results in a bit shift. Analogously, we can interpret a general multiplication as a bit shift by a fractional number of bits, followed by an (expensive) operation that has to “re-align” all valid bits to integer offsets. (ii) separation of the compressed representation into a bulk and a head.

Now consider the general case where $m(x_i)$ is not necessarily an integer power of two. Multiplying $m(x_i)$ to `self.compressed` grows the number of valid bits again by roughly $\log_2 m(x_i)$ bits (up to rounding effects). In fact, it is useful to think of the multiplication as a left bit shift by the non-integer amount of $\log_2 m(x_i)$ bits. In this picture, after the left-shift, the bits will *not* land at any positions that can actually store bits but they will instead land somewhere in-between, i.e., at a fractional alignment (see Figure 1 (i), which illustrates a bit shift by $\log_2 m(x_i) = 3.4$ bits). You can think of the multiplication with $m(x_i)$ as a fractional bit shift by $\log_2 m(x_i)$, followed by the task of “re-aligning” the bits to valid positions, which is expensive since it has to recalculate every valid bit.

To summarize,

- When we encode a symbol, we shift the existing bits on the buffer to the left by the information content of the symbol, and we then write the symbol into the generated space at the right (i.e., least significant) end of the buffer.
- When we decode a symbol, we read it off from the right end and then shift all bits to the right by the consumed information content.
- Integer bit shifts (aka normal bit shifts) are cheap, even on very large numbers, because we can implement them by growing or shrinking the buffer rather than actually shifting the bits.
- By multiplying with or dividing by an arbitrary number (which is not necessarily an integer power of two), we can achieve the equivalent of a “fractional bit shift”. However, such fractional bit shifts become increasingly expensive as the buffer grows and should therefore be avoided on large buffers.

Strategy. Motivated by the above observations, we will split `self.compressed` into two parts (see Figure 1 (ii)):

- a **bulk** part that always stores an integer amount of information content; and
- a **head** part that stores the remaining fractional amount of information content.

When we encode and decode symbols, we mostly operate on the **head** (which is therefore indeed the head of the stack). Importantly the **head** has a finite capacity, so that the cost for encoding and decoding operations on the **head** is bounded by a constant. Only when the **head** would overflow or underflow do we access the **bulk** to transfer some valid bits between the **head** and the **bulk**. These transfers will also have a constant (amortized) run-time cost because we only ever transfer an *integer* amount of bits.

Your task will be to figure out *when* and *how* exactly you have to transfer data between the **head** and the **bulk**.

- (a) Recall that we approximate probabilities by rational numbers, $P_{\text{ANS}}(X_i = x_i) = \frac{m(x_i)}{n}$ with $n = 2^{\text{precision}}$. What are the minimal and maximal information contents that a symbol with nonzero probability can have in this approximation?

Solution: The minimum information content corresponds to the highest possible probability, i.e., to $P_{\text{ANS}}(X_i = x_i) = 1$ and thus $m(x_i) = n$. In this case, the information content is $-\log_2 1 = 0$.

The maximum information content corresponds to the lowest possible probability. Within the fixed point precision $P_{\text{ANS}}(X_i = x_i) = \frac{m(x_i)}{2^{\text{precision}}}$ with integer $m(x_i)$, the smallest representable *nonzero* probability is $\frac{1}{2^{\text{precision}}} = 2^{-\text{precision}}$, and thus the corresponding information content is $-\log_2(2^{-\text{precision}}) = \text{precision}$.

Note that zero probability would correspond to *infinite* information content. This reflects the fact that an ANS coder *cannot* encode symbols with zero probability (it would lead to a division by zero in the **encode** method). This may not seem like a problem since symbols with zero probability shouldn't appear in the message. However, in practice, any probabilistic model of a data source will only be an approximation of the true data distribution, and a model might wrongfully assign zero probability even to symbols that can occur in the message. Thus, when approximating models with the fixed point representation P_{ANS} , one has to ensure that all symbols that might appear in the message have a nonzero probability. (The entropy coders in the **constriction** library ensure that all symbols within the alphabet are approximated with a nonzero probability, even if the provided floating point probability for a symbol is smaller than $2^{-\text{precision}}$.) ■

We will represent the **head** as a single integer and the **bulk** as a vector (in Python: a list) of integers. A popular (albeit not the only possible) strategy to manage the **bulk** and **head** is to uphold the following two invariants:

- **head** $< 2^{2 \times \text{precision}}$ (always); and
- **head** $\geq 2^{\text{precision}}$ if the **bulk** is not empty.

The jupyter notebook has a skeleton implementation of a class `StreamingAnsCoder` whose constructor initializes `self.head` and `self.bulk` in a way that trivially satisfies the above invariants. However, the `encode` and `decode` methods don't uphold these invariants yet. Instead, they currently just mirror the implementation from the lecture, except that they encode and decode to and from `self.head` instead of `self.compressed`.

- (b) Consider the `encode` method and assume that both of the above invariants hold at its entry. Show that the invariants will be violated at method exit *exactly* if, at method entry,

`(self.head >> self.precision) >= scaled_probabilities[symbol]`

Solution: For the analysis, denote $m(x_i) := \text{scaled_probabilities}[\text{symbol}]$ and $n := 2^{\text{precision}}$, and let h be the value of `self.head` at method entry and h' be the value of `self.head` at method exit. Thus, the above condition reads $\lfloor \frac{h}{n} \rfloor \geq m(x_i)$ and the `encode` method sets

$$h' = \left\lfloor \frac{h}{m(x_i)} \right\rfloor \times n + z$$

where $0 \leq z < n$.

We have to prove two things: (i) that h' violates at least one of the invariants if the condition $\lfloor \frac{h}{n} \rfloor \geq m(x_i)$ holds, and (ii) that h' satisfies both invariants if the condition $\lfloor \frac{h}{n} \rfloor \geq m(x_i)$ doesn't hold.

- (i) If the condition holds at method entry then we have

$$h \geq n \times \left\lfloor \frac{h}{n} \right\rfloor \geq n m(x_i)$$

and thus

$$h' = \left\lfloor \frac{h}{m(x_i)} \right\rfloor \times n + z \geq \left\lfloor \frac{n m(x_i)}{m(x_i)} \right\rfloor \times n = n^2 = 2^{2 \times \text{precision}}$$

which violates the first invariant.

- (ii) If the condition does not hold at method entry then we have $\lfloor \frac{h}{n} \rfloor \leq m(x_i) - 1$ and thus

$$h = n \times \frac{h}{n} < n \times \left(\left\lfloor \frac{h}{n} \right\rfloor + 1 \right) \leq n m(x_i).$$

This implies $\frac{h}{m(x_i)} < n$ and therefore $\lfloor \frac{h}{m(x_i)} \rfloor \leq n - 1$. Therefore, we have

$$h' = \left\lfloor \frac{h}{m(x_i)} \right\rfloor \times n + z \leq (n - 1) \times n + z = n^2 - n + z < n^2 = 2^{2 \times \text{precision}}$$

where we used $z < n$ in the last inequality. Thus, h' satisfies the first invariant.

The second invariant is satisfied by h' because it is satisfied by h by assumption, the current implementation doesn't mutate `bulk`, and $h' \geq h$ (because $z \geq h \% m(x_i) = h - m(x_i) \times \lfloor \frac{h}{m(x_i)} \rfloor$).

■

- (c) Modify the `encode` method so that it transfers `precision` bits from `self.head` to `self.bulk` if the case from part (b) arises at method entry. More precisely, it should pop the `precision` lowest significant bits off `self.head` and append a single item to `self.bulk` that represents these bits. Convince yourself that this operation will never transfer any (meaningless) leading zero bits to `self.bulk`. Then show that, with your modification, the method now upholds *both* of the above invariants (i.e., both invariants hold at method exit provided that they hold at method entry).

Solution: See accompanying jupyter notebook for the implementation.

The transfer from `self.head` to `self.bulk` never transfers any leading zero bits because it is only invoked if the condition from part (b) holds, which implies that `self.head >= (1 << self.precision)` (since `scaled_probabilities[symbol] > 0` by assumption). Note, however, that the high-order valid bits of `self.head` have slightly lower entropy than 1 bit due to Benford's Law.² Thus, streaming ANS has a small additional linear overhead. In practice, this overhead can be made negligibly small by either using large enough `precision` or by increasing the upper bound on `self.head` to some value larger than $2^{2 \times \text{precision}}$, so that transfers from `self.head` to `self.bulk` only ever transfer relatively low significant bits, whose entropy is very close to 1 bit.

To show that this modified method upholds the invariants, we only have to consider the case where the condition $\lfloor \frac{h}{n} \rfloor \geq m(x_i)$ holds at method entry. If it doesn't hold, then the proof from part (b) still applies. If the condition holds, then the modified method now sets

$$h' = \left\lfloor \frac{\lfloor h/n \rfloor}{m(x_i)} \right\rfloor \times n + z.$$

Since h satisfies the invariants, we have $h < n^2$ and thus $\lfloor h/n \rfloor < n$, which means $\lfloor \frac{\lfloor h/n \rfloor}{n} \rfloor = 0 < m(x_i)$. Thus, h' satisfies the first invariant by the same arguments as in item (ii) of the proof of part (b).

To prove that h' also upholds the second invariant, we start from the condition $\lfloor \frac{h}{n} \rfloor \geq m(x_i)$ and therefore obtain

$$h' = \left\lfloor \frac{\lfloor h/n \rfloor}{m(x_i)} \right\rfloor \times n + z \geq \left\lfloor \frac{m(x_i)}{m(x_i)} \right\rfloor \times n + z = n + z \geq n = 2^{\text{precision}}.$$

■

- (d) Now modify the `decode` method so that it becomes again the inverse of `encode` (i.e., if you start from any valid state and then call `encode` followed by `decode`, you always end up again in the original state). In detail, you will have to conditionally

²https://en.wikipedia.org/wiki/Benford%27s_law

transfer some data from `bulk` to `head`, and you should figure out both *where in the code* and *under which condition* this transfer must happen. Use the provided test to debug your implementation.

Solution: See accompanying jupyter notebook for the implementation. The `decode` method inverts all steps of the `encode` method in reverse order, which means that the conditional transfer from `self.bulk` to `self.head` happens at the *end* of the method. The condition for the transfer is

```
(self.head >> self.precision) == 0 and len(self.bulk) != 0
```

i.e., it checks whether the second invariant is violated. This works because the `encode` method upholds all invariants at method exit, but it *temporarily* violates the second invariant after the conditional transfer from `self.head` to `self.bulk` *exactly* if a transfer has occurred. Thus, checking whether the second invariant is violated allows us to detect whether a transfer occurred in the *encode* method and therefore has to be reversed in the *decode* method. ■

- (e) Now, add a simple method `get_compressed` that one can call at the end of encoding to get the full compressed representation, i.e., the concatenation of `self.bulk` and `self.head`. The method should split up `self.head` into two halves of `precision` bits each, and return a list of integers that is `self.bulk` followed the two halves of `self.head`, starting with the lower significant half. If the higher-significant half of `self.head` is zero then it should not be included in the returned value (according to the above invariants, this can only happen if `self.bulk` is empty). Finally, modify the constructor (`__init__`) so that it can accept an initial compressed representation (in the form returned by `get_compressed`). Debug your implementations again with the provided test.

Solution: See accompanying jupyter notebook for the implementation. ■

- (f) (*Bonus Question: Random Access*) Imagine you use your `StreamingAnsCoder` to compress some data, and you want the decoder to be able to quickly jump to certain (predefined) positions within the message without having to first decode everything that has been encoded on top of it in the stack. For example, you may want to allow the decoder to quickly jump to any symbol x_i with $i = 1000j$ for arbitrary integers j , and to then resume decoding from there. Thus, you'd package the compressed representation into some container format that also contains a *jump table*. What information should the jump table store for each desired target position so that the decoder can use this information to quickly jump there?

Solution: For each desired target position, the jump table has to store the offset into the compressed data and the value of `self.head` at this point. The latter can be regarded as a “fractional part” of the offset. ■




Figure 2: Stream coding with the naive ANS implementation from the lecture notes.

Problem 7.2: Streaming ANS II — Mondrian

This problem continues the discussion from Problem 7.1 but can be solved independently from the implementation of the algorithm.

Figure 2 illustrates how the naive ANS implementation from the lecture packs the encoded information content tightly into a bit string (as opposed to a symbol code, which would align each symbol to a bit boundary, thus introducing “gaps”). The illustrated compressed representation is the result of encoding the following 10 symbols, in this order:³

- a symbol x_1 with information content 1.5 bits;
- a symbol x_2 with information content 2.2 bits;
- a symbol x_3 with information content 1.4 bits;
- a symbol x_4 with information content 1.7 bits;
- a symbol x_5 with information content 1.9 bits;
- a symbol x_6 with information content 0.8 bits;
- a symbol x_7 with information content 1.1 bits;
- a symbol x_8 with information content 0.7 bits;
- a symbol x_9 with information content 1.6 bits; and
- a symbol x_{10} with information content 1.5 bits.

What would your `StreamingAnsCoder` from Problem 7.1 do if you used it to encode the same example sequence of symbols and then called `get_compressed()`? Draw a figure analogous to Figure 2 to sketch what the resulting compressed representation would look like (i.e., which parts of it correspond to which symbols). Assume `precision = 4`

³We’ll gloss over the fact that, in reality, ANS wouldn’t be able to represent these *precise* information contents since the corresponding probabilities $2^{-(\text{information content})}$ aren’t rational numbers.

(which would be an unreasonably low precision for real applications but suffices here for demonstration purpose).

You should find that some of the symbols get “split up” into two or even three non-neighboring parts, and that the very first symbol x_1 doesn’t get flushed from the **head** to the **bulk** until the very end. More precisely, you should end up with a compressed representation that is a sequence of four integers, each one carrying **precision** = 4 bits of information content where

- the first integer encodes x_3, x_4 , and a part of x_2 ;
- the second integer encodes x_6, x_7, x_8 , and a part of x_5 ;
- the third integer encodes x_9, x_{10} , another part of x_2 , and another part of x_5 ; and
- the fourth integer encodes x_1 and yet another part of x_2 .

Solution: The following figure illustrates how a streaming ANS coder with **precision** = 4 would encode this sequence of symbols. ■

